

LogisticRegression1

```
In [1]: import numpy as np  
import pandas as pd
```

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\C3_bot_detection_data - C3_bot_detection_data.csv")
df
```

Out[2]:

	User ID	Username	Tweet	Retweet Count	Mention Count	Follower Count	Verified	Bot Label	Location
0	132131	flong	Station activity person against natural majori...	85	1	2353	False	1	Adkin
1	289683	hinesstephanie	Authority research natural life material staff...	55	5	9617	True	0	Sander
2	779715	roberttran	Manage whose quickly especially foot none to g...	6	2	4363	True	0	Harrison
3	696168	pmason	Just cover eight opportunity strong policy which.	54	5	2242	True	1	Martinezl
4	704441	noah87	Animal sign six data good or.	26	3	8438	False	1	Camacho
...
49995	491196	uberg	Want but put card direction know miss former h...	64	0	9911	True	1	Kimberlybi
49996	739297	jessicamunoz	Provide whole maybe agree church respond most ...	18	5	9900	False	1	Greenl
49997	674475	lynn cunningham	Bring different everyone international capital...	43	3	6313	True	1	Deboral
49998	167081	richardthompson	Than about single generation itself seek sell ...	45	1	6343	False	0	Stephen

	User ID	Username	Tweet	Retweet Count	Mention Count	Follower Count	Verified	Bot Label	Loca
49999	311204	daniel29	Here morning class various room human true bec...	91	4	4006	False	0	Novaki

50000 rows × 11 columns



In [3]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   User ID         50000 non-null  int64
1   Username        50000 non-null  object
2   Tweet           50000 non-null  object
3   Retweet Count   50000 non-null  int64
4   Mention Count   50000 non-null  int64
5   Follower Count  50000 non-null  int64
6   Verified        50000 non-null  bool
7   Bot Label       50000 non-null  int64
8   Location        50000 non-null  object
9   Created At      50000 non-null  object
10  Hashtags        41659 non-null  object
dtypes: bool(1), int64(5), object(5)
memory usage: 3.9+ MB
```

```
In [5]: df1=df.fillna(0)
df1
```

Out[5]:

	User ID	Username	Tweet	Retweet Count	Mention Count	Follower Count	Verified	Bot Label	Loca
0	132131	flong	Station activity person against natural majori...	85	1	2353	False	1	Adkin
1	289683	hinesstephanie	Authority research natural life material staff...	55	5	9617	True	0	Sander
2	779715	roberttran	Manage whose quickly especially foot none to g...	6	2	4363	True	0	Harriso
3	696168	pmason	Just cover eight opportunity strong policy which.	54	5	2242	True	1	Martinezl
4	704441	noah87	Animal sign six data good or.	26	3	8438	False	1	Camacho
...
49995	491196	uberg	Want but put card direction know miss former h...	64	0	9911	True	1	Kimberlybi
49996	739297	jessicamunoz	Provide whole maybe agree church respond most ...	18	5	9900	False	1	Greenl
49997	674475	lynncunningham	Bring different everyone international capital...	43	3	6313	True	1	Deboral
49998	167081	richardthompson	Than about single generation itself seek sell ...	45	1	6343	False	0	Stephen

	User ID	Username	Tweet	Retweet Count	Mention Count	Follower Count	Verified	Bot Label	Loca
49999	311204	daniel29	Here morning class various room human true bec...	91	4	4006	False	0	Novaki

50000 rows × 11 columns

In [6]: df1.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User ID               50000 non-null  int64
1   Username              50000 non-null  object
2   Tweet                 50000 non-null  object
3   Retweet Count         50000 non-null  int64
4   Mention Count         50000 non-null  int64
5   Follower Count        50000 non-null  int64
6   Verified              50000 non-null  bool
7   Bot Label             50000 non-null  int64
8   Location              50000 non-null  object
9   Created At            50000 non-null  object
10  Hashtags              50000 non-null  object
dtypes: bool(1), int64(5), object(5)
memory usage: 3.9+ MB
```

In [16]: data=df1[['User ID', 'Retweet Count', 'Mention Count', 'Follower Count', 'Bot Label']]

In [17]: from sklearn.linear_model import LogisticRegression

In [18]: feature_matrix = data.iloc[:,0:5]
target_vector = data.iloc[:,1]

In [19]: feature_matrix.shape

Out[19]: (50000, 5)

In [20]: target_vector.shape

Out[20]: (50000,)

```
In [21]: from sklearn.preprocessing import StandardScaler
```

```
In [22]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [24]: logr=LogisticRegression()
logr.fit(fs,target_vector)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
 Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
 n_iter_i = _check_optimize_result(

```
Out[24]: LogisticRegression()
```

```
In [27]: observation=[[5,7,9,5,6]]
```

```
In [28]: prediction=logr.predict(observation)
print(prediction)
```

```
[100]
```

```
In [29]: logr.classes_
```

```
Out[29]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12,
                13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25,
                26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38,
                39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51,
                52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64,
                65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77,
                78, 79, 80, 81, 82, 83, 84, 85, 86, 87, 88, 89, 90,
                91, 92, 93, 94, 95, 96, 97, 98, 99, 100], dtype=int64)
```

```
In [30]: logr.predict_proba(observation)[0][0]
```

```
Out[30]: 7.962997325787375e-227
```

```
In [31]: logr.predict_proba(observation)[0][0]
```

```
Out[31]: 7.962997325787375e-227
```

LogisticRegression2

```
In [32]: import re
from sklearn.datasets import load_digits
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
```

```
In [33]: digits = load_digits()
digits
```


of hand written digits. 10 classes where (each class refers to a digit). (NIST preprocessing programs made available by NIST were used to extract normalized bitmaps of handwritten digits from a preprinted form. From a total of 43 people, 30 contributed to the training set and different 13 to the test set. 32 x32 bitmaps are divided into nonoverlapping blocks of 4x4 and the number of pixels are counted in each block. This generates an input matrix of 8x8 where each element is an integer in the range 0..16. This reduces dimensionality and gives invariance to small distortions.)

For info on NIST preprocessing routines, see M. D. Garris, J. L. Blue, G. T. Candela, D. L. Dimmick, J. Geist, P. J. Grother, S. A. Janet, and C. L. Wilson, NIST Form-Based Handprint Recognition System, NISTIR 5469, 1994.

.. topic:: References

- C. Kaynak (1995) Methods of Combining Multiple Classifiers and Their Applications to Handwritten Digit Recognition, MSc Thesis, Institute of Graduate Studies in Science and Engineering, Bogazici University.
- E. Alpaydin, C. Kaynak (1998) Cascading Classifiers, Kybernetika.
- Ken Tang and Ponuthurai N. Suganthan and Xi Yao and A. Kai Qin. Linear dimensionality reduction using relevance weighted LDA. School of Electrical and Electronic Engineering Nanyang Technological University. 2005.
- Claudio Gentile. A New Approximate Maximal Margin Classification Algorithm. NIPS. 2000.

```
In [37]: plt.figure(figsize=(20,4))
for index,(image,label) in enumerate(zip(digits.data[0:5],digits.target[0:5])):
    plt.subplot(1,5,index+1)
    plt.imshow(np.reshape(image,(8,8)))
    plt.title("number\n"%label,fontsize=15)
```



```
In [38]: x_train,x_test,y_train,y_test=train_test_split(digits.data,digits.target,test_size=0.2,random_state=0)
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(1257, 64)
(540, 64)
(1257,)
(540,)
```

```
In [39]: logre=LogisticRegression()
logre.fit(x_train,y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

Out[39]: LogisticRegression()

```
In [40]: print(logre.predict(x_test))
```

```
[1 4 4 9 9 2 9 0 2 0 1 0 0 6 3 5 7 2 2 1 9 5 8 7 0 2 6 3 6 5 6 8 1 5 3 4 3
 9 5 1 9 9 2 5 9 1 3 8 6 3 8 8 0 6 5 9 2 0 2 1 3 2 4 2 4 2 9 1 4 0 9 5 3 7
 1 0 0 8 3 2 4 3 2 9 8 6 6 3 5 0 2 1 0 5 2 5 6 2 0 2 1 1 0 8 7 4 9 0 3 3 7
 9 4 4 3 4 1 6 7 7 6 8 3 4 8 0 7 2 8 2 9 4 3 4 2 9 3 4 3 1 4 6 4 8 5 1 4 6
 6 0 4 3 7 3 4 7 9 4 6 1 8 9 7 3 4 5 7 4 7 7 0 9 9 4 6 2 0 3 2 0 6 2 1 0 6
 8 7 8 2 6 4 7 3 7 4 5 6 5 1 3 1 3 7 1 9 1 7 5 7 6 3 4 4 9 3 8 4 4 0 0 5 4
 1 0 1 8 7 4 9 3 2 2 0 4 4 8 5 5 1 4 5 6 0 1 0 3 3 7 7 1 0 3 2 0 0 4 6 1 8
 4 0 0 2 0 0 1 2 6 0 2 5 1 8 2 9 9 3 0 2 5 6 4 9 5 5 8 7 6 6 1 0 7 0 3 0 5
 3 9 6 0 5 6 8 5 4 7 7 4 5 5 4 1 6 0 1 7 5 2 6 6 9 5 8 5 7 5 2 6 1 4 5 1 5
 1 4 4 6 3 1 2 6 9 0 5 8 6 3 7 1 2 6 7 6 1 0 4 9 4 0 6 2 0 1 3 3 5 3 4 5 2
 5 9 8 5 8 7 7 9 0 1 4 5 0 6 1 6 9 6 3 5 3 7 0 9 0 9 1 0 9 8 5 3 2 4 2 4 0
 2 1 7 6 3 0 1 8 6 9 1 1 4 8 8 2 3 3 5 3 4 5 5 3 9 1 0 6 2 5 3 2 5 1 0 0 5
 7 8 5 8 7 8 5 8 9 9 1 8 5 7 3 2 2 9 1 6 3 9 7 1 0 0 5 2 8 0 0 8 0 0 6 6 1
 7 5 7 7 8 8 1 7 1 9 2 5 0 2 6 2 7 8 2 7 7 8 5 8 0 1 7 0 9 0 5 9 0 5 1 0 9
 8 2 5 5 9 3 9 0 3 6 4 9 3 9 4 9 9 0 0 8 3 3]
```

```
In [41]: print(logre.score(x_test,y_test))
```

```
0.9685185185185186
```